


Robustness — the distortions are structured signal and not encoder-specific

(A) Null model — real contexts vs a magnitude-matched random shift



(B) Multi-embedder — grounding holds across encoders (cross-encoder top-1 agreement 73%)

